

Adam Petrovic

Senior Software Engineer – Distributed Systems, Cloud Infrastructure & Reliability

Sydney, Australia

 adam@petrovic.com.au  [linkedin.com/in/adampetrovic](https://www.linkedin.com/in/adampetrovic)
 github.com/adampetrovic  adampetrovic.me

Summary

Senior software engineer with 14+ years building cloud platforms, object storage, and reliability systems, including 10+ years at Atlassian. Current work is on Atlassian's internal S3-compatible storage service: GCP migration, Kubernetes, storage analytics, incident response, and cost/reliability work. Earlier roles covered observability systems, a \$2M+ metrics migration, 180M metrics/minute pipelines, synthetic monitoring, and Jira Cloud SRE. Enjoys turning complexity and risk into RFCs, production code, and tooling adopted by other teams across the organisation.

Skills

- **Languages:** Kotlin, Go, Python, Java, TypeScript, JavaScript, SQL, Bash
- **Cloud & Infrastructure:** AWS, GCP, Kubernetes, Docker, Spinnaker, Flux, Helm, Kustomize, Talos Linux, CI/CD
- **Storage & Data Systems:** S3, GCS, DynamoDB, PostgreSQL, MySQL, Redis, sharding, backup/restore
- **Reliability & Observability:** SignalFx, Datadog, Prometheus, Grafana, Splunk, StatsD, Sentry, synthetic monitoring, SLOs
- **Data & Performance:** Kafka, Apache Flink, SQS, Databricks, Athena, Locust, high-volume telemetry pipelines

Experience

Atlassian (10+ Years)

Sydney, Australia

Senior Engineer, Transactional Data Platform (Object Store)

August 2023 – Present

- Designed Kotlin-based architecture for Atlassian's internal S3-compatible Object Store service, which manages 60PB+ of product data across the company, covering storage APIs, data residency, encryption, backup and restore, GCS integration, and multi-shard deployment across AWS and GCP.
- Led the GCP migration by re-architecting Object Store from a traditional microservice into a Kubernetes-native workload, modernising compute, credentials, data-integrity checks, and deployment readiness to deliver a critical company-wide OKR four weeks early.
- Helped shape Object Store's technical direction through RFCs and DACIs for GCS checksum handling and seekable compression, turning reliability, cost, and migration trade-offs into clear team decisions.
- Built S3 Inventory and DynamoDB export pipelines into Databricks, giving service owners and Object Store operators near real-time access to query 60PB+ of stored data for arbitrary questions about cost attribution, object verification, incident response, and production savings, turning multi-day data-integrity investigations into repeatable verification jobs with alerting and enabling \$25K/month in object-tag savings.
- Created a Bitbucket deployment diff pipeline that flags risky infrastructure changes before merge, adopted by multiple platform teams to catch blast-radius issues before production.
- Built repeatable Locust load tests for staging, letting engineers validate feature changes under production-like traffic, compare results across runs, and build confidence before deployment.

Engineering Manager, Observability

October 2020 – August 2023

- Led company-wide metrics migration from Datadog to SignalFx for Atlassian engineering, saving more than \$2M annually while keeping critical metrics workflows stable for thousands of engineers during transition.
- Managed observability engineers through vendor evaluation, roadmap planning, stakeholder communication, incident reviews, hiring, coaching, performance development, and operational trade-offs.
- Led delivery of an HTTP routing layer between metrics aggregators and upstream observability sinks, using Atlassian's open-source GoStatsD project to make the Datadog-to-SignalFx migration safer at scale.

Senior Engineer, Observability

June 2017 – October 2020

- Built and operated Atlassian's metrics infrastructure at 180M datapoints/minute, supporting production monitoring, alerting, and incident diagnosis across cloud services and internal platforms.
- Developed Pollinator, a Go-based synthetic monitoring and post-deployment verification platform used across regions to catch reliability and performance regressions in release flows before customers did.
- Extended GoStatsD with Prometheus-style histograms and moved telemetry toward Kafka/Flink, improving percentile metrics, stream-processing headroom, and SLO diagnosis for service teams.

Senior Site Reliability Engineer, Jira Cloud

March 2015 – June 2017

- Led a cross-continental SRE group delivering Jira Cloud's European infrastructure ahead of a public launch milestone, supporting enterprise customers, regional growth, and product commitments.
- Reworked Jira Cloud site imports and introduced Wheel of Misfortune training, reducing failed customer imports, on-call toil, and incident-response gaps through rehearsed failure scenarios.

Freelancer.com (3.5 Years)

Sydney, Australia

Software Engineer

August 2011 – February 2015

- Built platform services for authentication, messaging, notifications, email, analytics, and fraud systems in a global marketplace using Python, Go, Redis, RabbitMQ, Thrift, and async workers.
- Helped move critical paths away from PHP monolith patterns into service-oriented APIs and async processing, improving request latency, failure isolation, and delivery speed for core marketplace flows.

Education

The University of Sydney

Sydney, Australia

Bachelor of Information Technology, Computer Science

2008 – 2012

- Dean's List with a distinction average.