# Adam Petrovic

Senior Software Engineer – Distributed Systems, Cloud Infrastructure & Reliability

Sydney, Australia

adam@petrovic.com.au linkedin.com/in/adampetrovic

github.com/adampetrovic adampetrovic.me


## Summary


Senior software engineer with 14+ years building cloud platforms, object storage, and reliability systems, including 10+ years at Atlassian. Current work is on Atlassian's internal S3-compatible storage service: GCP migration, Kubernetes, storage analytics, incident response, and cost/reliability work. Earlier roles covered observability systems, a $2M+ metrics migration, 180M metrics/minute pipelines, synthetic monitoring, and Jira Cloud SRE. Enjoys turning complexity and risk into RFCs, production code, and tooling adopted by other teams across the organisation.


## Skills

- **Languages:** Kotlin, Go, Python, Java, TypeScript, JavaScript, SQL, Bash
- **Cloud & Infrastructure:** AWS, GCP, Kubernetes, Docker, Spinnaker, Flux, Helm, Kustomize, Talos Linux, CI/CD
- **Storage & Data Systems:** S3, GCS, DynamoDB, PostgreSQL, MySQL, Redis, sharding, backup/restore
- **Reliability & Observability:** SignalFx, Datadog, Prometheus, Grafana, Splunk, StatsD, Sentry, synthetic monitoring, SLOs
- **Data & Performance:** Kafka, Apache Flink, SQS, Databricks, Athena, Locust, high-volume telemetry pipelines

## Experience

**Atlassian** (11+ Years) — **Sydney, Australia**

*Senior Engineer, Transactional Data Platform (Object Store)* — *August 2023 – Present*

- Designed Kotlin-based architecture for Atlassian's internal S3-compatible Object Store service, which manages 60PB+ of product data across the company, covering storage APIs, data residency, encryption, backup and restore, GCS integration, and multi-shard deployment across AWS and GCP.
- Led the GCP migration by re-architecting Object Store from a traditional microservice into a Kubernetes-native workload, modernising compute, credentials, data-integrity checks, and deployment readiness to deliver a critical company-wide OKR four weeks early.
- Helped shape Object Store's technical direction through RFCs and DACIs for GCS checksum handling and seekable compression, turning reliability, cost, and migration trade-offs into clear team decisions.
- Built S3 Inventory and DynamoDB export pipelines into Databricks, giving service owners and Object Store operators near real-time access to query 60PB+ of stored data for arbitrary questions about cost attribution, object verification, incident response, and production savings, turning multi-day data-integrity investigations into repeatable verification jobs with alerting and enabling $25K/month in object-tag savings.
- Created a Bitbucket deployment diff pipeline that flags risky infrastructure changes before merge, adopted by multiple platform teams to catch blast-radius issues before production.
- Built repeatable Locust load tests for staging, letting engineers validate feature changes under production-like traffic, compare results across runs, and build confidence before deployment.

*Engineering Manager, Observability* — *October 2020 – August 2023*

- Led company-wide metrics migration from Datadog to SignalFx for Atlassian engineering, saving more than $2M annually while keeping critical metrics workflows stable for thousands of engineers during transition.
- Managed observability engineers through vendor evaluation, roadmap planning, stakeholder communication, incident reviews, hiring, coaching, performance development, and operational trade-offs.
- Led delivery of an HTTP routing layer between metrics aggregators and upstream observability sinks, using Atlassian's open-source GoStatsD project to make the Datadog-to-SignalFx migration safer at scale.

*Senior Engineer, Observability* *June 2017 – October 2020*

- Built and operated Atlassian's metrics infrastructure at 180M datapoints/minute, supporting production monitoring, alerting, and incident diagnosis across cloud services and internal platforms.
- Developed Pollinator, a Go-based synthetic monitoring and post-deployment verification platform used across regions to catch reliability and performance regressions in release flows before customers did.
- Extended GoStatsD with Prometheus-style histograms and moved telemetry toward Kafka/Flink, improving percentile metrics, stream-processing headroom, and SLO diagnosis for service teams.

*Senior Site Reliability Engineer, Jira Cloud* *March 2015 – June 2017*

- Led a cross-continental SRE group delivering Jira Cloud's European infrastructure ahead of a public launch milestone, supporting enterprise customers, regional growth, and product commitments.
- Reworked Jira Cloud site imports and introduced Wheel of Misfortune training, reducing failed customer imports, on-call toil, and incident-response gaps through rehearsed failure scenarios.

**Freelancer.com** (3.5 Years) **Sydney, Australia**
*Software Engineer* *August 2011 – February 2015*

- Built platform services for authentication, messaging, notifications, email, analytics, and fraud systems in a global marketplace using Python, Go, Redis, RabbitMQ, Thrift, and async workers.
- Helped move critical paths away from PHP monolith patterns into service-oriented APIs and async processing, improving request latency, failure isolation, and delivery speed for core marketplace flows.

## Education

**The University of Sydney** **Sydney, Australia**
*Bachelor of Information Technology, Computer Science* *2008 – 2012*

- Dean's List with a distinction average.